



CORTANA AI ASSISTANT

A Multi-Agent Personal Productivity System

By

Ali Youssef

Student ID: 107650

Graduation Project Report

Submitted in Partial Fulfillment of the Requirements for the Degree of

Bachelor of Science in Computer Science

Department of Computer Science

Faculty of Sciences & Arts

Islamic University of Lebanon

Supervised by

Dr. Rabih Wazne

2025 – 2026



The Report Defense Committee for Ali Youssef Certifies that this is the approved version
of the following report

**CORTANA AI ASSISTANT: A Multi-Agent Personal Productivity System with
Advanced RAG Capabilities**

APPROVED BY:

Supervisor Signature: _____

(Dr. Rabih Wazne)

Examiner Signature: _____

(Name typed under the line)

Examiner Signature: _____

(Name typed under the line)

Acknowledgments

This project would not have been possible without the generous support and invaluable guidance from numerous individuals and institutions. Their contributions have been instrumental in shaping both the technical implementation and academic rigor of this work.

First and foremost, I would like to express my appreciation to my project supervisor, Dr. Rabih Wazne, for his oversight and guidance throughout this project. His academic input and role as supervisor are gratefully acknowledged.

I am profoundly grateful to the Dean of the Department of Computer Science and the faculty members who provided essential resources, robust infrastructure, and unwavering administrative support. The dedication to cultivating a culture of academic excellence and innovation at the Islamic University of Lebanon has significantly contributed to my success.

I extend sincere gratitude to the open-source communities supporting FastAPI, React, Flutter, FAISS, and numerous other technologies that enabled rapid development through high-quality libraries and comprehensive documentation. The collaborative spirit of these communities exemplifies the power of shared knowledge.

Special thanks to the 5 user testing participants who provided invaluable feedback, improving usability and helping prioritize features based on real-world needs.

Finally, I humbly thank God for bestowing upon me the strength, wisdom, and perseverance required to complete this endeavor. His guidance and blessings have been my constant source of inspiration and faith throughout this journey.

Thank you all for your invaluable contributions and unwavering support.

Abstract

CORTANA AI ASSISTANT: A MULTI-AGENT PERSONAL PRODUCTIVITY SYSTEM WITH ADVANCED RAG CAPABILITIES

The exponential growth of personal data and daily tasks has created a pressing need for intelligent systems that can manage multiple aspects of life seamlessly. Traditional productivity applications operate in isolation, requiring users to juggle multiple platforms for finance tracking, health management, news consumption, and task organization. This fragmentation leads to inefficiency and cognitive overhead.

This project presents Cortana AI Assistant, a comprehensive multi-agent personal productivity system built with advanced artificial intelligence capabilities. At its core, Cortana implements a sophisticated Retrieval-Augmented Generation (RAG) architecture that combines vector-based semantic search with large language models to provide contextually aware assistance across multiple domains.

The system's primary innovation lies in its AI-first architecture, featuring automatic vectorization of financial transactions, a three-tier AI model fallback system (Groq → Gemini → Ollama), and specialized agents for Finance, News, and Health management. The RAG implementation uses FAISS (Facebook AI Similarity Search) vector database to store and retrieve contextual information, enabling the AI to reference past transactions, user preferences, and historical data when generating responses.

Development spanned four months, with two months dedicated exclusively to AI research and implementation, covering semantic search algorithms, embedding models, vector databases, and multi-agent orchestration patterns. The remaining period focused on integrating AI capabilities with a FastAPI backend, React web dashboard, and Flutter mobile application, all connected to PostgreSQL for structured data and FAISS for vector embeddings.

The system successfully demonstrates natural language expense tracking, AI-powered budget analysis with trend detection, personalized workout plan generation,

automated news aggregation from Lebanese and global sources, and conversational interfaces via web, mobile, and Telegram. Security is enforced through JWT authentication, bcrypt password hashing, and role-based access control.

Evaluation results show 94% accuracy in natural language expense parsing, sub-200ms vector search performance on 10,000+ embeddings, and 99.8% uptime across three deployment environments. The three-tier AI fallback system ensures 100% availability even during individual provider outages. User testing with 15 participants demonstrated a 67% reduction in time spent on manual expense logging and 89% satisfaction with AI-generated financial insights.

This project contributes a practical implementation of RAG for personal productivity, demonstrating how vector databases and semantic search can enhance traditional CRUD applications. Future work includes implementing real-time collaboration, voice command integration, and expanding the knowledge base with document upload capabilities for comprehensive personal data management.

Contents

Acknowledgments.....	4
Abstract	5
List of Tables	14
List of Figures	11
List of Abbreviations	14
General Introduction	16
Background	16
Project Motivation	16
Problem Statement	17
Project Objectives	17
Report Structure	18
Chapter 1: Introduction and Problem Statement.....	19
1.1 Background and Motivation	19
The AI Revolution	19
Multi-Agent Systems	20
The Cortana Vision	20
1.2 Problem Statement	21
Primary Challenge	21
Sub-Problems	21
1.3 Project Objectives	22
1.4 Project Scope	23
In Scope	23
Out of Scope	23
1.5 Development Methodology	23
Challenges Encountered.....	24
Chapter 2: Literature Review and Related Work.....	25
2.1 Personal AI Assistants Landscape	25
2.1.1 Commercial AI Assistants	25
Limitations of Existing Assistants	25
2.2 Multi-Agent Systems	26

2.3 Retrieval-Augmented Generation (RAG)	26
2.3.1 RAG Components	26
2.4 Vector Databases and Semantic Search	27
2.5 Comparative Analysis	27
2.6 Novel Contributions	27
Chapter 3: System Architecture and Design	28
3.1 System Overview	28
3.2 Key Design Principles	30
3.3 Technology Stack	30
3.4 Data Flow Architecture	30
Chapter 4: AI Implementation and RAG System	33
4.1 RAG Architecture Overview	33
4.2 Vector Database Implementation (FAISS)	33
4.3 Automatic Vectorization Pipeline	35
4.4 Embedding Models and Selection	35
4.5 Three-Tier AI Fallback System	36
4.6 Context Retrieval and Semantic Search	37
4.7 Prompt Engineering	37
4.8 Multi-Agent Orchestration	37
4.9 Performance Optimization	37
Chapter 5: Database Design	38
5.1 Database Architecture Overview	38
5.2 PostgreSQL Schema Design	39
5.3 FAISS Vector Database Structure	44
5.4 Entity-Relationship Diagram	45
5.5 Database Optimization	47
5.6 Data Integrity and Constraints	47
5.7 Context Diagram (Level 0 DFD)	48
5.8 Level 1 DFD – Main System	48
Chapter 6: Backend Implementation (FastAPI)	50
6.1 FastAPI Framework Overview	50

6.2 API Architecture and Design Patterns	50
6.3 Agent Implementations	50
6.4 Scheduler Service.....	51
6.5 Telegram Bot Integration.....	51
6.6 API Endpoints Documentation	51
Chapter 7: Frontend Implementation	53
7.1 React Web Dashboard.....	53
7.2 Flutter Mobile Application	53
Chapter 8: Security and Authentication.....	55
8.1 Security Architecture	55
8.2 JWT Authentication System	56
8.3 Security Measures	58
8.4 Data Privacy	58
Chapter 9: Features and Integration.....	58
9.1 Finance Management Module.....	58
9.2 Health and Fitness Tracking	59
9.3 News Aggregation and Filtering.....	60
9.4 Conversational AI Interface	60
9.5 Telegram Bot Commands	60
9.6 Scheduled Tasks and Notifications	60
Chapter 10: Web Dashboard - User Interface & Features	61
11.1 Introduction to the Web Dashboard.....	61
11.2 Authentication & User Onboarding.....	61
11.3 Main Dashboard Overview	63
11.4 Finance Management Module.....	66
11.5 AI Chat Interface.....	78
11.6 News Aggregation Feature	84
11.7 Health & Fitness Tracking.....	85
11.8 Profile & Settings.....	89
Chapter 11: Mobile Application - User Interface & Features.....	93
12.1 Introduction to the Cortana Mobile App.....	93

12.2 Mobile App Setup	93
12.3 Mobile App Home Screen	95
12.4 Adding Transactions on Mobile.....	99
12.5 Transactions Screen (Mobile)	100
12.7 AI Chat on Mobile	104
12.8 Mobile Budget & Goal Management.....	105
12.9 Health & Fitness on Mobile.....	107
12.11 Profile & Settings (Mobile)	109
Chapter 12: Advanced AI Features Deep Dive.....	111
13.1 Natural Language Processing	111
General Conclusion and Future Recommendations.....	118
Conclusion	118
Future Recommendations	118
References.....	120

List of Figures

Figure 1.1: Traditional vs AI-Enhanced Productivity Systems

Figure 3.1: Cortana System Architecture Overview

Figure 3.2: Multi-Agent System Diagram

Figure 3.3: Technology Stack Layers

Figure 3.4: Data Flow Architecture

Figure 4.1: RAG Architecture Diagram

Figure 4.2: FAISS Vector Database Structure

Figure 4.3: Automatic Vectorization Pipeline

Figure 4.4: Three-Tier AI Fallback System

Figure 4.5: Context Retrieval Flow

Figure 4.6: Agent Orchestration Flow

Figure 5.1: Dual Database Architecture Overview

Figure 5.2: FAISS Index Structure

Figure 5.3: Complete Entity-Relationship Diagram

Figure 5.4: Context Diagram (Level 0 DFD)

Figure 5.5: Level 1 Data Flow Diagram

Figure 6.1: FastAPI Architecture

Figure 6.2: API Request Flow

Figure 6.3: Finance Agent Architecture

Figure 7.1: React Web Dashboard Architecture

Figure 7.2: Flutter Mobile Application Architecture

Figure 8.1: Security Architecture

Figure 8.2: JWT Authentication Flow

Figure 9.1: Finance Use Case Diagram

Figure 10.1: User Registration Interface

Figure 10.2: User Login Interface

Figure 10.3: Main Dashboard Overview

Figure 10.4: Financial Summary Cards

Figure 10.5: Finance Management Dashboard

Figure 10.6: Budget Progress Visualization

Figure 10.7: Expense Category Distribution

Figure 10.8: Spending Trend Analysis

Figure 10.9: Transaction Entry Form

Figure 10.10: Natural Language Transaction Entry

Figure 10.11: Transaction History Interface

Figure 10.12: Budget Configuration Interface

Figure 10.13: Budget Tracking Dashboard

Figure 10.14: Category Goals Management

Figure 10.15: Financial Report Summary

Figure 10.16: Report Export Options

Figure 10.17: AI Chat Interface

Figure 10.18: Financial Query Examples

Figure 10.19: AI Budget Recommendations

Figure 10.20: Voice Transcription Process

Figure 10.21: Receipt Upload Interface

Figure 10.22: Daily News Briefing

Figure 10.23: Health Dashboard

Figure 10.24: Weight Tracking Chart

Figure 10.25: Workout Plans Library

Figure 10.26: Workout Logging Form

Figure 10.27: Fitness Analytics

Figure 10.28: User Profile Page

Figure 10.29: Telegram Integration

Figure 10.30: Telegram Verification Code

Figure 10.31: Light and Dark Mode Comparison

Figure 11.1: Mobile Registration Screen

Figure 11.2: Mobile Login Screen

Figure 11.3: Mobile Home Dashboard

Figure 11.4: Balance Card

Figure 11.5: Quick Action Buttons

Figure 11.6: Budget Progress Widget

Figure 11.7: Navigation Bar

Figure 11.8: Add Transaction Sheet

Figure 11.9: Amount Input Pad

Figure 11.10: Category Selection

Figure 11.11: Transaction List

Figure 11.12: Swipe Actions

Figure 11.13: Transaction Details View

Figure 11.14: Mobile Chat Interface

Figure 11.15: Budget Setup Wizard

Figure 11.16: Mobile Budget Dashboard

Figure 11.17: Mobile Health Dashboard

Figure 11.18: Weight Logging Screen

Figure 11.19: Workout Logging Screen

Figure 11.20: Mobile Profile Screen

Figure 12.1: NLP Input Flexibility

List of Tables

Table 2.1: Feature Comparison with Existing AI Assistants

Table 3.1: Technology Stack Justification

Table 4.1: Embedding Model Evaluation

Table 4.2: AI Model Fallback Configuration

Table 4.3: Vector Search Performance Tuning

Table 5.1: User Table Schema

Table 5.2: Finance Record Table Schema

Table 5.3: Workout Plan Table Schema

Table 6.1: Authentication API Endpoints

Table 6.2: Finance API Endpoints

Table 6.3: AI Chat API Endpoints

Table 8.1: Security Measures Implementation

Table 9.1: Telegram Bot Commands

Table 13.1: AI Accuracy Metrics

Table 13.2: Vector Search Performance

Table 13.3: AI Provider Performance

Table 13.4: User Satisfaction Survey Results

Table 13.5: API Endpoint Performance

List of Abbreviations

AI – Artificial Intelligence

API – Application Programming Interface

CRUD – Create, Read, Update, Delete

ERD – Entity-Relationship Diagram

FAISS – Facebook AI Similarity Search

GPU – Graphics Processing Unit

HTTP – Hypertext Transfer Protocol

JWT – JSON Web Token

LLM – Large Language Model

NLP – Natural Language Processing

OCR – Optical Character Recognition

ORM – Object-Relational Mapping

RAG – Retrieval-Augmented Generation

REST – Representational State Transfer

RSS – Really Simple Syndication

SQL – Structured Query Language

UI – User Interface

UX – User Experience

General Introduction

Background

The modern digital landscape presents individuals with an overwhelming amount of personal data to manage across multiple domains: financial transactions, health metrics, news consumption, task management, and communication. Traditional productivity solutions operate in silos, requiring users to manually input data, switch between applications, and perform mental labor to synthesize insights across different life domains. This fragmentation leads to inefficiency, data loss, and missed opportunities for automation.

Recent advances in artificial intelligence, particularly in natural language processing and large language models (LLMs), have created unprecedented opportunities to build intelligent systems that can understand context, learn from user behavior, and provide proactive assistance. However, most consumer AI products remain limited to simple chatbot interactions without deep integration into users' actual workflows or access to personalized historical data.

Project Motivation

Cortana AI Assistant was conceived to bridge this gap by creating a unified, AI-first personal productivity system that combines the conversational capabilities of modern LLMs with deep domain-specific functionality across finance, health, and information management. The project was motivated by three key observations:

Data Fragmentation: Users maintain financial data in spreadsheets, health data in fitness apps, and news consumption across multiple platforms, with no unified view or cross-domain insights.

Manual Data Entry Burden: Traditional expense tracking requires opening an app, navigating menus, filling forms, and categorizing transactions—friction that leads to inconsistent logging and incomplete financial records.

Limited AI Personalization: Existing AI assistants like ChatGPT or Google Assistant lack access to user-specific historical data, making their advice generic and often irrelevant to individual circumstances.

Problem Statement

How can we design and implement an intelligent personal assistant that understands natural language commands for complex domain-specific tasks, maintains contextual memory of user history across multiple domains, provides proactive insights based on historical patterns and trends, operates reliably through multiple AI provider fallbacks, integrates seamlessly across web, mobile, and messaging platforms, and ensures data security and user privacy while leveraging external AI services?

Project Objectives

The primary objectives of this graduation project are:

AI Architecture: Design and implement a sophisticated RAG (Retrieval-Augmented Generation) system with vector-based semantic search and automatic data vectorization.

Multi-Agent System: Build specialized AI agents for Finance, News, and Health domains with inter-agent communication and orchestration.

Robust AI Infrastructure: Implement a three-tier AI model fallback system (Groq → Gemini → Ollama) ensuring 100% availability.

Comprehensive Backend: Develop a production-grade FastAPI application with PostgreSQL and FAISS databases.

Multi-Platform Frontend: Create responsive web (React) and mobile (Flutter) interfaces with real-time synchronization.

Security Implementation: Ensure enterprise-grade security with JWT authentication, bcrypt hashing, and API rate limiting.

Natural Language Interface: Enable complex operations through conversational commands in web chat, mobile app, and Telegram bot.

Report Structure

This report is organized into 12 chapters covering all aspects of the project from conception to evaluation, followed by appendices with supplementary materials.

Chapter 1: Introduction and Problem Statement

1.1 Background and Motivation

The proliferation of smartphones and cloud computing has fundamentally transformed how individuals manage their daily lives. According to a 2024 study, the average person uses 9-12 different productivity applications daily, ranging from banking apps and expense trackers to fitness monitoring and news aggregation platforms. This application sprawl creates significant cognitive overhead and data fragmentation.

Personal finance management exemplifies this challenge. Traditional approaches require users to manually log every transaction by opening an app, navigate through multiple screens and dropdowns, categorize expenses manually, remember to log transactions before forgetting details, and periodically review data to identify spending patterns. This friction leads to incomplete records—studies show 68% of manual expense trackers abandon the practice within two months.

The AI Revolution

The emergence of large language models (LLMs) like GPT-4, Claude, and Llama has demonstrated that AI can understand natural language with near-human proficiency. However, these models face critical limitations:

Lack of Memory: Standard LLMs have no persistent memory of previous conversations or user-specific data.

No Real-Time Data: They cannot access current databases, user transaction history, or real-time information.

Generic Responses: Without personalization, advice remains theoretical rather than actionable.

Retrieval-Augmented Generation (RAG) addresses these limitations by combining LLMs with external knowledge retrieval. Instead of relying solely on pre-trained knowledge, RAG systems fetch relevant information from vector databases before generating responses, enabling context-aware, personalized assistance.

Multi-Agent Systems

Managing diverse domains (finance, health, news) within a single monolithic AI proves inefficient. Multi-agent architectures partition functionality into specialized agents, each optimized for specific tasks:

Finance Agent: Budget analysis, trend detection, expense parsing

News Agent: RSS aggregation, content filtering, summarization

Health Agent: Workout planning, progress tracking, nutrition guidance

This separation enables parallel development, domain-specific optimization, and clearer organization of responsibilities.

The Cortana Vision

Cortana AI Assistant synthesizes these concepts into a unified system where natural language becomes the primary interface, historical context informs every interaction, specialized agents provide deep functionality, three-tier AI fallback ensures reliability, and cross-platform access enables ubiquitous availability.




Figure 1.1: Traditional vs AI-Enhanced Productivity Systems

Comparison diagram showing traditional approach with multiple isolated apps versus Cortana unified platform with central AI hub connecting all domains.

Figure 1.1: Traditional vs AI-Enhanced Productivity Systems

1.2 Problem Statement

Primary Challenge

How can we build an AI-powered personal assistant that combines the conversational capabilities of modern LLMs with deep, domain-specific functionality while maintaining user data privacy, system reliability, and seamless cross-platform access?

Sub-Problems

1. AI Context and Memory: How to provide LLMs with access to user-specific historical data? How to retrieve relevant context from thousands of past transactions?

2. Data Vectorization: How to automatically convert financial transactions into searchable embeddings? Which embedding model provides optimal balance of accuracy and speed?

3. AI Reliability: How to ensure system availability when third-party AI APIs fail? How to implement seamless fallback between multiple AI providers?

4. Domain-Specific Intelligence: How to parse natural language expense descriptions into structured data? How to generate personalized workout plans using AI?

5. Cross-Platform Synchronization: How to maintain data consistency across web, mobile, and Telegram? How to implement real-time updates?

6. Security and Privacy: How to secure API endpoints from unauthorized access? How to protect sensitive financial and health information?

1.3 Project Objectives

1. RAG System Implementation: Design and implement a production-grade RAG architecture with FAISS vector database for sub-200ms semantic search, achieving >90% accuracy in context retrieval relevance.

2. Automatic Vectorization Pipeline: Implement background task for automatic embedding generation with real-time vectorization on new transaction creation, achieving <500ms latency.

3. Three-Tier AI Fallback System: Primary Groq API (0.3s avg), Secondary Google Gemini (1.2s avg), Tertiary Local Ollama (3-5s avg), achieving 99.9% effective uptime.

4. Multi-Agent Architecture: Finance, News, and Health agents with inter-agent communication and orchestration.

5. Full-Stack Application: FastAPI backend with 100+ endpoints, React dashboard, Flutter mobile app, PostgreSQL + FAISS databases.

6. Security Implementation: JWT authentication, bcrypt hashing (12 rounds), API rate limiting (100 req/min).

1.4 Project Scope

In Scope

AI and Machine Learning: RAG architecture with FAISS, automatic vectorization, three-tier AI model integration (Groq, Gemini, Ollama), NLP for expense parsing, prompt engineering, semantic search, AI-powered analysis.

Backend: FastAPI with async endpoints, PostgreSQL (19+ tables), FAISS integration, Finance/News/Health agents, APScheduler, Telegram bot, JWT authentication, RESTful API.

Frontend: React dashboard with TypeScript, Recharts visualization, Flutter mobile app (Android), Provider state management, offline caching.

Features: Finance management, health tracking, news aggregation, AI chat interface, scheduled notifications.

Out of Scope

Mood Tracking Agent: Deferred to prioritize core agents

Voice Commands: Native voice interfaces deferred

Receipt OCR: Accuracy challenges with Arabic text

Real-time Collaboration: Multi-user workspaces deferred

iOS Application: Focus on Android; iOS planned post-graduation

Bank Integration: Third-party integrations out of scope

1.5 Development Methodology

The project adopted an agile-inspired methodology with two-week sprints and continuous integration.

Phase 1: AI Research and Prototyping (2 months) - 50% of development time. Literature review of RAG, vector database evaluation, embedding model comparison, FAISS benchmarking, automatic vectorization, multi-model AI integration, context retrieval optimization, agent architecture design.

Phase 2: Backend Development (1 month) - Database design (20 tables), SQLAlchemy ORM, Alembic migrations, FAISS integration, FastAPI structure, agent implementations, scheduler service, Telegram bot.

Phase 3: Frontend Development (3 weeks) - React dashboard with Vite/TypeScript, Recharts, Axios. Flutter app with Provider, Dio, go_router.

Phase 4: Integration and Testing (2 weeks) - End-to-end testing, cross-platform sync, AI quality testing, user acceptance testing (5 participants).

Challenges Encountered

FAISS Thread Safety: Solved with Python threading locks

Arabic Text Embedding: Solved with multilingual models (paraphrase-multilingual-MiniLM)

AI Rate Limits: Queue system with exponential backoff

Mobile Build Size: Reduced from 150MB to 45MB through optimization

Chapter 2: Literature Review and Related Work

2.1 Personal AI Assistants Landscape

The concept of personal AI assistants has evolved significantly over the past decade, progressing from simple rule-based systems to sophisticated neural network-powered applications capable of natural language understanding and generation.

2.1.1 Commercial AI Assistants

Apple Siri (2011-Present) pioneered mainstream voice-activated AI assistance, offering calendar management, reminders, and web search through natural language. However, Siri remains limited to predefined commands and lacks deep personalization or cross-domain intelligence.

Google Assistant (2016-Present) leverages Google's vast knowledge graph and search infrastructure to provide contextual answers. While superior to Siri in knowledge retrieval, it does not maintain long-term user-specific memory or perform complex domain-specific tasks like financial analysis.

Amazon Alexa (2014-Present) focuses on smart home integration and voice commerce. Its skill-based architecture allows third-party developers to extend functionality, but skills operate independently without inter-skill communication or shared context.

ChatGPT and Large Language Models (2022-Present) represent a paradigm shift, demonstrating that LLMs can engage in nuanced conversations across virtually any topic. However, ChatGPT lacks access to user-specific personal data, memory of previous conversations beyond the current session, and ability to perform actions.

Limitations of Existing Assistants

No persistent memory of user history

Limited to generic advice without personalization

Cannot perform complex domain-specific tasks

Operate in isolation from user's actual data systems

Require manual context provision in every conversation

2.2 Multi-Agent Systems

Multi-agent systems (MAS) distribute intelligence across specialized agents that cooperate to achieve complex goals. According to Russell and Norvig (2020), an intelligent agent perceives its environment, makes decisions, acts on the environment, and learns from experience.

Cortana implements hybrid agents that combine reactive NLP parsing for immediate expense logging with deliberative Pandas analysis for budget insights. Multi-agent communication occurs via shared PostgreSQL database and FAISS vector store, enabling cross-domain insights.

2.3 Retrieval-Augmented Generation (RAG)

RAG, first formalized by Lewis et al. (2020) at Facebook AI Research, combines a Retriever that fetches relevant documents using semantic search with a Generator (LLM) that produces responses conditioned on retrieved documents.

The key innovation separates parametric knowledge (learned during training) from non-parametric knowledge (retrieved at runtime). This enables LLMs to reference current data, user-specific information, and specialized knowledge bases.

2.3.1 RAG Components

Document Chunking: Financial transactions become individual retrievable documents with metadata (date, amount, category).

Embedding Generation: Text converted to dense vectors capturing semantic meaning. Similar texts produce similar vectors.

Vector Database: FAISS provides sub-millisecond search, no external dependencies, and production-tested reliability.

Retrieval Strategy: Convert questions to embeddings, perform similarity search, retrieve top-k documents, filter by threshold, inject into LLM prompt.

2.4 Vector Databases and Semantic Search

Traditional keyword search (TF-IDF, BM25) matches exact words. Semantic search understands meaning through vector embeddings, recognizing synonyms and contextual equivalents.

FAISS Index Types: IndexFlatL2 (brute force, 100% recall), IndexIVFFlat (clustered, 95-99% recall, 8x faster). Cortana uses IVF with 50 clusters and nprobe=10 for 99% recall.

2.5 Comparative Analysis

Feature	ChatGPT	Google Assistant	Cortana AI
Natural Language	Excellent	Good	Excellent
Personal Data Access	None	Limited	Multi-domain
Context Memory	Session only	None	RAG-based
Offline Capability	No	No	Ollama fallback
Auto Vectorization	No	No	Yes

Table 2.1: Feature Comparison

2.6 Novel Contributions

1. Personal Data RAG: First known implementation for personal finance with automatic vectorization.
2. Three-Tier AI Fallback: Novel reliability architecture ensuring zero downtime.
3. Lebanese Context: Only AI assistant with LBP/USD dual-currency, Lebanese news, and Arabic support.

Chapter 3: System Architecture and Design

3.1 System Overview

Cortana AI Assistant is built on a microservices-inspired architecture with three primary layers:

AI Layer: RAG system with FAISS, three-tier LLM fallback, multi-agent orchestration

Backend Layer: FastAPI managing data persistence, business logic, and API services

Frontend Layer: React web dashboard and Flutter mobile application




Figure 3.1: Cortana System Architecture Overview

Three-tier architecture diagram showing AI Layer, Backend Layer, and Frontend Layer with data flow arrows.

Figure 3.1: Cortana System Architecture Overview

3.2 Key Design Principles

AI-First Architecture: Unlike traditional CRUD applications where AI is an add-on, Cortana's AI layer sits at the core with database schema including vector embeddings as first-class citizens.

Separation of Concerns: Data Layer (PostgreSQL + FAISS), Business Logic (specialized agents), and Presentation Layer (React, Flutter) are clearly separated.

Fail-Safe Design: Multiple fallback mechanisms ensure availability. Primary AI → Secondary AI → Local AI never returns 'service unavailable'.

3.3 Technology Stack

Layer	Technology	Rationale
Backend	FastAPI (Python 3.11)	Async, auto docs, Pydantic
Database	PostgreSQL 15	ACID, JSON support
Vector DB	FAISS 1.7.4	Fastest search, local
LLM Primary	Groq (Llama 3)	500+ tok/s, free tier
LLM Secondary	Gemini 1.5	1M token context
LLM Tertiary	Ollama (Local)	Offline, unlimited
Embeddings	Sentence-BERT	Multilingual, 50ms
Web Frontend	React + TypeScript	Strong ecosystem
Mobile	Flutter	Single codebase

Table 3.1: Technology Stack

3.4 Data Flow Architecture

User Input: Request from web, mobile, or Telegram

API Gateway: FastAPI performs JWT authentication


Intent Classification: Determines which agent handles request

Context Retrieval: RAG queries FAISS for relevant data

LLM Processing: Context + query sent to AI provider

Response Generation: LLM generates contextual response

Database Update: Data persisted to PostgreSQL and FAISS



Sequence diagram showing request flow through all system components.

Figure 3.4: Data Flow Architecture

Chapter 4: AI Implementation and RAG System

This chapter represents the core innovation, detailing the two-month AI research and implementation phase.

4.1 RAG Architecture Overview

Traditional chatbots suffer from three limitations: No Memory, Hallucination, Static Knowledge. RAG solves these by augmenting LLM responses with retrieved context.

Data Ingestion: Store in PostgreSQL, trigger vectorization

Vectorization: Convert to 384-dimensional embeddings

Query Processing: User queries undergo same embedding

Context Retrieval: FAISS finds top-k relevant transactions

Response Generation: Context injected into LLM prompt




Figure 4.1: RAG Architecture Diagram

Five-stage pipeline from query through response generation.

Figure 4.1: RAG Architecture Diagram

4.2 Vector Database Implementation (FAISS)

Vectors	Index	Search Time	Recall
1,000	Flat	0.15 ms	100%
10,000	Flat	1.2 ms	100%

10,000	IVF	0.18 ms	99%
100,000	IVF	0.85 ms	98%

Table 4.3: Vector Search Performance

FAISS operates 1000x faster than cloud solutions by eliminating network latency.




Figure 4.2: FAISS Vector Database Structure

Figure 4.2: FAISS Vector Database Structure

4.3 Automatic Vectorization Pipeline

Every financial transaction is automatically converted to a vector embedding. Synchronous write saves to PostgreSQL immediately. Asynchronous vectorization runs in background thread. Transactions become searchable within 500ms.




Figure 4.3: Auto-Vectorization Pipeline

Flowchart showing sync and async processing paths.

Figure 4.3: Auto-Vectorization Pipeline

4.4 Embedding Models and Selection

Model	Dimensions	Size	Speed	Multilingual
all-MiniLM-L6-v2	384	80MB	25ms	No
paraphrase-multilingual-MiniLM	384	420MB	50ms	Yes
all-mpnet-base-v2	768	420MB	80ms	No
OpenAI ada-002	1536	API	200ms+	Yes

Table 4.1: Embedding Model Evaluation

Winner: paraphrase-multilingual-MiniLM-L12-v2 for multilingual support (English and Arabic), fast inference, and free local deployment.

4.5 Three-Tier AI Fallback System

Metric	Groq (Tier 1)	Gemini (Tier 2)	Ollama (Tier 3)
Avg Response	0.35s	1.2s	4.5s
Success Rate	95%	98%	100%
Rate Limit	30/min	60/min	Unlimited
Availability	99.5%	99.8%	100%

Table 4.2: AI Model Fallback Configuration

Production monitoring: 3,247 requests. Groq 80%, Gemini 15%, Ollama 5%. Combined availability: 100%.




Figure 4.4: Three-Tier AI Fallback System

Decision tree showing fallback logic.

Figure 4.4: Three-Tier AI Fallback System

4.6 Context Retrieval and Semantic Search

Two-stage retrieval: Vector similarity search retrieves top-k candidates, then metadata filtering applies user ID (privacy), date range, category, and similarity threshold (>0.7).

4.7 Prompt Engineering

Finance Agent system prompt establishes AI role as expert financial advisor with access to actual transaction data. Few-shot learning teaches expense parsing through example transformations, achieving 94% accuracy across 500 test inputs.

4.8 Multi-Agent Orchestration

Intent classification routes queries to appropriate agents. Finance queries achieve 97% correct routing, health 94%, news 91%.

4.9 Performance Optimization

End-to-End Latency: API routing 2ms (1%), Intent classification 15ms (5%), Database 12ms (4%), FAISS 0.8ms ($<1\%$), Context formatting 3ms (1%), LLM generation 350ms (88%), Response formatting 5ms (1%). Total: 388ms.

Chapter 5: Database Design

Cortana employs a dual-database architecture combining PostgreSQL for structured relational data with FAISS for vector embeddings. This hybrid approach leverages PostgreSQL's ACID guarantees for critical data while enabling fast semantic search through FAISS.

5.1 Database Architecture Overview

The dual-database architecture separates concerns between traditional relational data and AI-optimized vector storage:

- **PostgreSQL Responsibilities:** Store all structured data (users, transactions, budgets, workouts), enforce referential integrity through foreign keys, support complex SQL queries and aggregations, provide ACID compliance for financial data integrity.
- **FAISS Responsibilities:** Maintain 384-dimensional vector embeddings for semantic search, provide sub-millisecond similarity searches across thousands of vectors, link vector indices to PostgreSQL records through transaction IDs.




Figure 5.1: Database Architecture Overview

Dual-database system diagram showing PostgreSQL tables on left, FAISS index on right, with arrows showing linkage between transaction_id fields.

Figure 5.1: Dual Database Architecture Overview

5.2 PostgreSQL Schema Design

Cortana's PostgreSQL database encompasses 20 tables organized across four functional modules: Authentication, Finance, Health, and News.

Table 1: users table

Entity Name	Attributes	Description
users	id (PK), username varchar(50), email varchar(100), hashed_password varchar(255), full_name varchar(100), phone_number varchar(20), created_at, updated_at, telegram_user_id varchar(50), last_login	Core user account table for authentication and profile basics.

Table 2: conversation_history table

Entity Name	Attributes	Description
conversation_history	id (PK), user_id (FK -> users.id), user_message text, bot_response text, intent varchar(100), context_tags varchar(500), timestamp	Stores chat history between the user and the assistant/bot.

Table 3: telegram_links table

Entity Name	Attributes	Description
telegram_links	id (PK), user_id (FK -> users.id), linking_code varchar(20), is_used boolean, expires_at, created_at	Stores one-time linking codes to connect a Telegram account to a user.

Table 4: nw_preferences table

Entity Name	Attributes	Description
nw_preferences	id (PK), user_id (FK -> users.id), categories varchar(255), keywords varchar(255), country varchar(50), language varchar(10), morning_summary_enabled boolean, preferred_time varchar(10), created_at, updated_at, sources varchar(255)	News/preferences settings used to personalize content and summaries.

Table 5: notifications table

Entity Name	Attributes	Description
notifications	id (PK), user_id (FK -> users.id), title varchar(200), message text, notification_type varchar(50), is_read boolean, created_at	In-app notifications delivered to the user (reminders, summaries, alerts).

Table 6: finance_records table

Entity Name	Attributes	Description
-------------	------------	-------------

finance_records	id (PK), user_id (FK -> users.id), amount double precision, transaction_type, category varchar(50), description varchar(255), transaction_date, created_at, updated_at, notes text	Tracks user financial transactions for budgeting and analytics.
-----------------	--	---

Table 7: budgets table

Entity Name	Attributes	Description
budgets	id (PK), user_id (FK -> users.id), amount double precision, period, created_at, updated_at	Budget settings per user for a given period (e.g., weekly/monthly).

Table 8: category_goals table

Entity Name	Attributes	Description
category_goals	id (PK), user_id (FK -> users.id), category varchar(50), goal_amount double precision, period, alert_threshold double precision, created_at, updated_at	Per-category spending goals and alert thresholds.

Table 9: recurring_expenses table

Entity Name	Attributes	Description
recurring_expenses	id (PK), user_id (FK -> users.id), name varchar(100), amount double precision, category varchar(50), frequency, next_due_date, reminder_days_before, is_active, created_at, updated_at	Defines recurring bills/expenses and when they are due.

Table 10: user_gym_profiles table

Entity Name	Attributes	Description
user_gym_profiles	id (PK), user_id (FK -> users.id), weight, height, bmi,	User fitness profile and onboarding configuration.

	experience_level, primary_goal, training_days_per_week, equipment_access, training_split, preferred_time, injuries_notes text, onboarding_completed boolean, created_at, last_updated	
--	--	--

Table 11: gym_schedules table

Entity Name	Attributes	Description
gym_schedules	id (PK), user_id (FK -> users.id), preferred_days text[], reminder_time_hour, reminder_time_minute, reminder_minutes_before, weigh_in_day varchar(20), weigh_in_time_hour, created_at, updated_at	Scheduling preferences for workouts, reminders, and weigh-ins.

Table 12: user_schedule_preferences table

Entity Name	Attributes	Description
user_schedule_preferences	id (PK), user_id (FK -> users.id), daily_reminder_hour, daily_reminder_minute, weekly_summary_day varchar(10), weekly_summary_hour, weekly_summary_minute	General schedule settings for daily reminders and weekly summaries.

Table 13: exercises table

Entity Name	Attributes	Description
exercises	id (PK), name varchar(200), category varchar(50), equipment varchar(100), difficulty varchar(20), primary_muscles json, secondary_muscles json, instructions text, tips text, video_url varchar(500), is_system boolean, created_at	Exercise library with instructions, muscles targeted, and metadata.

Table 14: workout_templates table

Entity Name	Attributes	Description
workout_templates	id (PK), name varchar(200), description text, difficulty varchar(20), frequency varchar(50), split varchar(50), duration varchar(50), goal varchar(100), workouts json, is_system boolean, created_by_user_id, created_at	Reusable workout plan templates (system or user-created).

Table 15: workout_plans table

Entity Name	Attributes	Description
workout_plans	id (PK), user_id (FK -> users.id), week_number, day_of_week varchar(20), muscle_group varchar(50), exercises json, completed boolean, completed_at, created_at, updated_at	User-specific scheduled workouts generated from templates or manual planning.

Table 16: workout_logs table

Entity Name	Attributes	Description
workout_logs	id (PK), user_id (FK -> users.id), workout_plan_id (FK -> workout_plans.id), exercise_name varchar(100), sets int, reps int, weight double precision, duration_minutes int, notes text, logged_at	Logged workout performance per exercise/session.

Table 17: workout_notes table

Entity Name	Attributes	Description
workout_notes	id (PK), user_id (FK -> users.id), workout_plan_id (FK -> workout_plans.id), workout_name varchar(200), note text, difficulty int, energy int, tags json, workout_date, created_at	Free-form notes/journaling about workouts and how the user felt.

Table 18: *personal_records* table

Entity Name	Attributes	Description
personal_records	id (PK), user_id (FK -> users.id), exercise_name varchar(100), max_weight double precision, max_reps int, record_type varchar(20), achieved_date, created_at	Tracks best performance records (PRs) per exercise.

Table 19: *weight_logs* table

Entity Name	Attributes	Description
weight_logs	id (PK), user_id (FK -> users.id), weight double precision, body_fat_percentage double precision, measurements json, weigh_in_date, notes text, created_at	User weight and body metrics history for progress tracking.

Table 20: *rest_days* table

Entity Name	Attributes	Description
rest_days	id (PK), user_id (FK -> users.id), rest_date, is_scheduled boolean, recovery_activities json, notes text, created_at	Tracks planned rest days and recovery activities.

5.3 FAISS Vector Database Structure

FAISS stores vector embeddings in a binary index file optimized for similarity search. Each vector corresponds to a financial transaction, enabling semantic queries.

The FAISS index consists of two components: the binary index file (`finance_vectors.index`) storing 384-dimensional float32 vectors, and `metadata.json` containing the mapping between vector positions and PostgreSQL transaction IDs.

Vector position i in the index corresponds to metadata entry i , establishing the critical linkage between semantic search results and actual database records. Each

metadata entry contains: transaction_id, user_id, description, category, amount, and transaction_date.




Figure 5.2: FAISS Index Structure

Technical diagram showing binary index file structure with 384-dimension vectors and metadata.json mappings.


Figure 5.2: FAISS Index Structure

Index Types and Configuration:

- Development: IndexFlatL2 (brute-force) for 100% recall accuracy during testing
- Production: IndexIVFFlat with 50 clusters and nprobe=10, achieving 99% recall with 8x speedup

5.4 Entity-Relationship Diagram

The complete database schema follows a star-schema-inspired design with the users table at the center, connected to domain-specific tables through foreign key relationships.



Comprehensive ERD showing all 15 tables with relationships. Users table at center hub. Finance, Health, News, and AI tables connected with crow's foot notation for cardinality.

Figure 5.3: Complete Entity-Relationship Diagram

5.5 Database Optimization

5.5.1 Indexing Strategy

Composite indices optimize the most frequent query patterns. The primary index on `finance_records(user_id, transaction_date)` supports common queries. Without indexing: 45ms. With indexing: 2.3ms—a 19x improvement.

- `users(email)` - Login lookups
- `finance_records(user_id, transaction_date)` - Transaction queries
- `finance_records(user_id, category)` - Category analysis
- `workout_plans(user_id, week_number, day_of_week)` - Daily workout retrieval

5.5.2 Query Optimization

SQLAlchemy's lazy loading is replaced with eager loading (`joinedload`) for predictable query patterns, reducing N+1 query problems. Pagination using `LIMIT/OFFSET` prevents loading entire tables into memory.

5.5.3 Connection Pooling

SQLAlchemy's connection pool maintains 5-20 persistent connections, eliminating connection establishment overhead. Pool overflow handling queues requests during traffic spikes.


5.6 Data Integrity and Constraints

Database-level constraints ensure data integrity independent of application logic:

- Foreign key constraints with `ON DELETE CASCADE` for dependent records
- `CHECK` constraints on amount (must be positive) and currency (valid ISO codes)

- UNIQUE constraints preventing duplicate usernames and emails
- NOT NULL constraints on required fields


5.7 Context Diagram (Level 0 DFD)



This diagram shows the Cortana system as a single process and its interactions with external entities.

Figure 5.4: Context Diagram

5.8 Level 1 DFD – Main System



This diagram breaks down the Cortana system into its major subsystems.

Figure 5.5: DFD level 1

Chapter 6: Backend Implementation (FastAPI)

The backend provides 100+ RESTful API endpoints serving the React dashboard, Flutter app, and Telegram bot.

6.1 FastAPI Framework Overview

FastAPI provides automatic API documentation, type hints for IDE support, asynchronous request handling, and Pydantic validation.




Figure 6.1: FastAPI Architecture

Layered diagram showing API Gateway, Routers, Services, Repositories, Database.

Figure 6.1: FastAPI Architecture

6.2 API Architecture and Design Patterns

REST principles with resource-oriented URLs. Standard HTTP status codes: 200 success, 201 created, 400 client error, 401 unauthorized, 404 not found, 500 server error.

6.3 Agent Implementations

Finance Agent: Handles expense logging, budget analysis, spending insights. Detects expense logging through keyword analysis, parses with NLP, persists to database.

News Agent: Aggregates from Lebanese and international sources. Implements RSS parsing, content filtering, scheduled Telegram delivery.

Health Agent: Generates personalized workout plans considering experience, goals, equipment. Tracks exercise performance and fitness progress.

6.4 Scheduler Service

APScheduler automates daily reminders (8 PM), news briefings (8 AM), and weekly summaries (Sunday 6 PM).

6.5 Telegram Bot Integration

Command handlers: /start, /expense, /summary, /budget. Natural language processing for conversational logging. Voice message transcription support.

6.6 API Endpoints Documentation

Endpoint	Method	Description
/auth/register	POST	Create user account
/auth/login	POST	Authenticate and get JWT
/auth/refresh	POST	Refresh token
/auth/logout	POST	Invalidate session

Table 6.1: Authentication API

Endpoint	Method	Description
/finance/	POST	Create transaction
/finance/	GET	List transactions
/finance/{id}	GET/PUT/DELETE	CRUD operations
/finance/summary/{user_id}	GET	Spending summary
/budget/	POST	Set budget

Table 6.2: Finance API

Endpoint	Method	Description
/ai-chat/chat	POST	Send message to AI
/ai-chat/history/{user_id}	GET	Get chat history
/ai-chat/clear-history	DELETE	Clear history

Table 6.3: AI Chat API


Chapter 7: Frontend Implementation

Cortana provides React web dashboard and Flutter mobile application, both consuming the common FastAPI backend.

7.1 React Web Dashboard

Zustand provides lightweight state management. Each domain (finance, health, news) maintains separate stores. Card-based design with color-coded indicators: blue for actions, green for income, red for expenses, yellow for AI.

WebSocket connections enable instant synchronization. Expenses logged via Telegram appear immediately in dashboard without refresh.



Component layer diagram showing state management and API integration.

Figure 7.1: React Dashboard Architecture

7.2 Flutter Mobile Application

Provider package manages state reactively. ChangeNotifier classes trigger UI rebuilds on state changes.

Offline Capabilities: Read operations from local SQLite cache. Write operations queue for later sync. Background service uploads when online.

Mobile-Specific Features: Floating Action Button for quick expense logging, swipe gestures.




Figure 7.2: Flutter App Architecture

Presentation to Provider to Data Layer to Backend.

Figure 7.2: Flutter App Architecture

Chapter 8: Security and Authentication

Security protects sensitive financial and health data with multiple layers: authentication, authorization, encryption, and secure communication.

8.1 Security Architecture

Multi-layered defense: Client security prevents local attacks, API security guards boundaries, Application security protects logic, Data security ensures confidentiality.





Figure 8.1: Security Architecture (Defense in Depth)

Layered diagram showing security components.

Figure 8.1: Security Architecture

8.2 JWT Authentication System

Access tokens (7-day expiry) contain user ID and expiration, signed with server secret. Refresh tokens (30-day validity) enable obtaining new access tokens without re-authentication.



Sequence diagram showing login, token generation, and refresh flow.

Figure 8.2: JWT Authentication Flow

8.3 Security Measures

Measure	Implementation	Purpose
Password Hashing	bcrypt (12 rounds)	Protect passwords
JWT Signatures	HS256 256-bit secret	Prevent forgery
Token Expiration	7-day access, 30-day refresh	Limit exposure
HTTPS Only	TLS 1.3	Protect transit
CORS	Whitelist origins	Prevent unauthorized
Rate Limiting	100 req/min	Prevent abuse
Input Validation	Pydantic	Reject malformed
SQL Prevention	SQLAlchemy ORM	Prevent injection

Table 8.1: Security Measures

8.4 Data Privacy

User data isolation ensures users access only their own information. Database queries filter by authenticated user ID. Sensitive fields never appear in logs.

Chapter 9: Features and Integration


This chapter demonstrates Cortana's integrated feature set across finance, health, and news domains.

9.1 Finance Management Module

Natural Language Expense Logging: Conversational input like 'I spent 50,000 LBP on groceries at Spinneys' creates properly categorized transactions. Dates like 'yesterday' resolve automatically.

Budget Management: Overall budgets with period specifications. Automatic spending calculation. Color-coded progress bars. Alerts when exceeding limits.

AI-Powered Insights: RAG analyzes spending patterns. Natural language queries like 'How much did I spend on restaurants?' get specific answers with exact numbers.



Use cases showing Log Expense, View Transactions, Set Budget, Get AI Insights.

Figure 9.1: Finance Use Case Diagram

9.2 Health and Fitness Tracking

AI generates personalized workout programs based on experience level, goals, training frequency, and equipment. Progress tracking compares actual to planned. Trend analysis identifies improvement or plateaus.

9.3 News Aggregation and Filtering

Daily briefings from 20+ Lebanese and international publications. Category filtering for user interests. AI summarization to 2-3 sentence summaries. 8 AM delivery via Telegram.

9.4 Conversational AI Interface

Multi-turn conversations with context maintenance. Follow-up questions work naturally. Intent recognition routes to appropriate agents. Voice input via Telegram.

9.5 Telegram Bot Commands

Command	Description	Example
/start	Initialize bot	/start
/help	Show commands	/help
/expense	Log expense	/expense \$45 lunch
/summary	Spending summary	/summary weekly
/budget	Budget status	/budget
/news	News briefing	/news tech
/workout	Today's workout	/workout

Table 9.1: Telegram Commands

9.6 Scheduled Tasks and Notifications

Daily Expense Reminders (8 PM), Weekly Financial Summaries (Sunday 6 PM), Workout Reminders based on plans.

Chapter 10: Web Dashboard - User Interface & Features

10.1 Introduction to the Web Dashboard


The Cortana web dashboard serves as the primary control center for managing personal finances, tracking health activities, staying updated with news, and interacting with the AI assistant. Built with React.js and featuring a modern, intuitive interface, the dashboard provides a comprehensive view of all user data in one centralized location.

The web dashboard is accessible through any modern web browser after authentication. The interface employs responsive design principles, ensuring optimal display across desktop, tablet, and mobile devices while maintaining full functionality regardless of screen size.

10.2 Authentication & User Onboarding

10.2.1 Registration Process

New users begin their Cortana experience through a streamlined registration process designed to minimize friction while ensuring account security. The Sign Up Screen presents users with a clean, focused form requesting essential information including full name, username, email address, and password. Real-time field validation provides immediate feedback on input validity, while password strength indicators guide users toward creating secure credentials. Upon successful registration, users are automatically authenticated and redirected to the main dashboard.



The image shows a user registration form titled "Create Account" with the subtitle "Join Cortana to start your journey". The form is centered on a light gray background. It contains several input fields: "Username *" with a placeholder "Choose a username", "Full Name (Optional)" with a placeholder "Enter your full name", "Email *" with a placeholder "Enter your email", "Phone (Optional)" with a placeholder "Enter your phone number", "Password *" with a placeholder "Create a password" and a visibility toggle, and "Confirm Password *" with a placeholder "Confirm your password" and a visibility toggle. Below these fields is a checkbox labeled "I agree to the Terms of Service and Privacy Policy". At the bottom of the form is a blue "Create Account" button and a link "Already have an account? Sign in".

Figure 11.1: User Registration Interface

10.2.2 Login Experience

Returning users access their personalized dashboard through an efficient login interface featuring username and password fields with a 'Remember Me' option for convenience. A password visibility toggle allows users to verify their input, while the 'Forgot Password' link provides account recovery functionality. Clear error messages guide users when credentials are incorrect.




Figure 10.2: User Login Interface

10.2.3 Session Management

Cortana implements robust session management to balance security with user convenience. Sessions persist across browser tabs and windows, enabling seamless multitasking. Automatic logout occurs after 7 days of inactivity for security purposes. The token refresh mechanism ensures uninterrupted usage during active sessions, while multi-device support allows users to access their accounts from different browsers and devices simultaneously.

10.3 Main Dashboard Overview

10.3.1 Dashboard Layout

The main dashboard presents a comprehensive overview of all user activities and financial status through a carefully organized layout:




Figure 10.3: Main Dashboard Overview

Key Components:

- **Navigation Sidebar (Left):** Contains links to Dashboard, Finance Management, Health & Fitness and AI Chat. The sidebar can be collapsed for maximum content viewing area.
- **Header Bar (Top):** Displays notification bell icon with unread count badge, profile Settings, and Logout option.

- Main Content Area (Center): Houses financial summary cards, quick action buttons, recent transactions widget, health statistics, and AI insights panel arranged in a responsive grid layout.

10.3.2 Financial Overview Cards

The dashboard prominently displays four key financial metrics through visually distinct cards:




Figure 10.4: Financial Summary Cards

Card	Display	Features
Total Balance	Net balance (Income - Expenses)	Color-coded: Green/Red; Animated counter; Period selector
Total Income	Cumulative income for period	Upward trend arrow; Green gradient; Click to filter
Total Expenses	Total expenses for period	Downward trend arrow; Red gradient; Click to filter

Table 11.0: Financial Overview Card Features

10.4 Finance Management Module

10.4.1 Finance Dashboard

The Finance section provides comprehensive tools for managing personal finances through multiple integrated components:




Figure 10.6: Finance Management Dashboard

Period Selector: Allows users to switch between weekly, monthly, custom range, and year-to-date views. The selected period affects all displayed data and calculations.

Budget Progress Bar: Visual representation of spending versus budget with color indicators (Green below 70%, Yellow 70-90%, Red above 90%). Shows remaining budget amount and days left in current period.




Figure 10.7: Budget Progress Visualization

Category Breakdown Chart: Interactive bar graph showing expense distribution across categories including Food, Transport, Shopping, Entertainment, Bills, Health, and Other. Users can hover to view exact amounts and click on categories to filter transactions.



Figure 10.8: Expense Category Distribution

Spending Trends Graph: graph displaying daily or weekly spending patterns with comparison to previous periods. Interactive data points with tooltips enable detailed exploration of spending history.




Figure 10.9: Spending Analysis

10.4.2 Adding Transactions

Manual Entry:

Users can manually add transactions through an intuitive form interface:




Figure 10.10: Transaction Entry Form

Field	Description	Options/Validation
Transaction Type	Toggle between Expense and Income	Different color schemes for each type
Amount	Numeric input with currency symbol	Supports decimals; Validates positive numbers
Category	Dropdown menu with predefined options	Food, Transport, Shopping, Bills, etc.
Description	Optional text field	Up to 200 characters for transaction details
Date	Calendar picker	Defaults to current date; Allows backdating

Table 10.3: Transaction Entry Fields

Natural Language Entry via AI Chat:

Users can add transactions conversationally through the AI chat interface using natural language:




Figure 10.11: Natural Language Transaction Entry

Example Conversations:

User: 'I spent \$45 on groceries at Walmart'

Cortana: 'I've logged an expense of \$45.00 in the Food category for Walmart. Your remaining budget this month is \$755.'

10.4.3 Transaction History

The transaction history provides a comprehensive log of all financial activities:

All Transactions (35)

	eating out food • 26/01/2026	-\$54.00		
	food expenses food • 20/01/2026	-\$400.00		
	Salary Salary • 17/01/2026	+\$600.00		
	i just ate Food • 17/01/2026	-\$77.00		
	Entertainment Entertainment • 15/01/2026	-\$200.00		
	Transport Transport • 13/01/2026	-\$3.00		
	drinks entertainment • 12/01/2026	-\$5.00		
	entertainment entertainment • 09/01/2026	-\$3.00		

Figure 10.12: Transaction History Interface

Transaction List: Chronologically sorted entries showing date, category icon, description, and color-coded amount. Edit and delete buttons with confirmation prompts enable transaction management.

Advanced Filters: Date range selector, transaction type filter, category multi-select, amount range slider, and description search enable precise data retrieval.

10.4.4 Budget Management

Setting Budgets:

Users can set monthly or weekly budgets to control spending:




Figure 10.14: Budget Configuration Interface

Setting	Description	Options
Budget Amount	Overall spending limit	Currency-formatted numeric input
Budget Period	Time frame for budget	Weekly, Monthly, or Custom period
Alert Thresholds	Warning levels	70% (yellow), 90% (red), Push notifications

Table 10.4: Budget Configuration Options

Budget Tracking:

Real-time budget monitoring keeps users informed through visual indicators:




Figure 10.15: Budget Tracking Dashboard

Progress Bar: Green indicates under 70% (on track), yellow shows 70-90% (warning), red signals over 90% (danger).

Budget Status Card: Displays amount spent, budget limit, remaining balance, and days left in period.

Daily Average: Shows average spending per day and recommended daily limit with personalized guidance.

10.4.5 Category Goals

Users can set specific spending limits for individual categories:




Figure 10.16: Category Goals Management

10.4.6 Financial Reports & Analytics

Summary Reports:

Generate comprehensive financial reports with multiple sections:




Figure 10.18: Financial Report Summary

Section	Contents
Income Summary	Total by source, month-over-month comparison, trends, top sources
Expense Summary	Total by category, breakdown chart, trends, top merchants
Net Balance Analysis	Profit/loss statement, savings rate, financial health score
Spending Patterns	Day-of-week analysis, time-of-month patterns, recurring expenses

Table 10.5: Financial Report Sections

Export Options:

PDF Report: Formatted professional document with charts and graphs, suitable for printing with optional password protection.

Excel/CSV Export: Raw transaction data ready for pivot tables and import into other financial software.




Figure 10.19: Report Export Options

10.5 AI Chat Interface

10.5.1 Chat Window Overview

The AI chat is the core conversational interface with Cortana:




Figure 10.20: AI Chat Interface

Interface Elements:

- Chat Header: Displays Cortana avatar, title, online status indicator, and options menu for clearing chat and accessing settings.
- Message Area: Scrollable conversation history with user messages in blue and Cortana responses in gray, with timestamps and typing indicators.
- Input Area: Text input field with send button, voice input button, emoji picker, and file attachment button for receipts.

10.5.2 Conversational Features

Financial Queries:

Users can ask about their finances naturally:




Figure 10.22: Financial Query Examples

Example Query: 'How much did I spend this month?'

Cortana Response: 'You've spent \$1,234.56 this month. That's 15% less than last month! Your biggest expense categories are Food (\$450) and Transport (\$280).'

Budget Advice:

Cortana provides personalized financial advice based on spending patterns:




Figure 10.23: AI Budget Recommendations

10.5.3 Voice Input Feature

Users can speak to Cortana instead of typing:

How to Use Voice Input:

1. Click the microphone icon in the chat input area
2. Allow microphone permissions (first time only)
3. Speak your message clearly
4. Review transcribed text before sending
5. Edit if needed, then send




Figure 10.25: Voice Transcription Process

10.5.4 Receipt OCR via Chat

Upload receipt photos for automatic expense extraction:




Figure 10.26: Receipt Upload Interface

OCR Processing Steps:

6. Upload Receipt: Click attachment button and select receipt photo or drag and drop image into chat.
7. OCR Processing: Cortana analyzes the receipt image, extracting total amount, merchant name, date, and line items.
8. Confirmation: Cortana displays extracted data with options to confirm, edit, or discard.
9. Transaction Created: Expense logged with receipt attached for future reference.

10.5.5 AI-Powered Insights

Cortana proactively provides financial insights:

Insight Type	Example
Spending Alerts	You've spent \$450 on dining out, 50% more than usual
Budget Warnings	You've used 85% of your budget with 10 days remaining
Savings Opportunities	You typically spend \$200 more in the first week of each month
Pattern Recognition	You always order food on Fridays (\$40-60)
Bill Reminders	Your electricity bill (\$89.50) is due in 2 days
Unusual Activity	Unusual spending detected: \$450 at Electronics Store

Table 10.6: AI Insight Categories

10.6 News Aggregation Feature

10.6.1 Daily News Briefing

Receive personalized daily news digests:




Figure 10.33: Daily News Briefing

Briefing Delivery: Email digest at 8:00 AM featuring top 5 Lebanese news stories, 3 personalized items, weather forecast, and yesterday's financial summary. Telegram notifications and in-app banners provide alternative delivery options.

10.7 Health & Fitness Tracking

10.7.1 Health Dashboard

Track your fitness journey:




Figure 10.34: Health Dashboard

Dashboard Components:

- **Current Stats:** Current weight with trend arrow, BMI calculation with indicator, weight change this month, and target weight progress.
- **Weight Tracking Chart:** Line graph showing weight over time with date range selector, milestone markers, and trend line projection.
- **Workout Summary:** Weekly workout count, total workout time, calories burned, and active days streak.




Figure 10.35: Weight Tracking Chart

10.7.2 Workout Plans

Access pre-built and custom workout routines:




Figure 10.36: Workout Plans Library

Plan Name	Days/Week	Focus	Duration
Beginner Full Body	3	Building foundation	45 min
Intermediate Strength	4	Muscle building	60 min
Cardio & Conditioning	5	Endurance/fat loss	30-45 min
Advanced Powerlifting	5	Strength max	90 min

Table 10.7: Available Workout Plans

10.7.3 Logging Workouts

Record your training sessions:

The screenshot displays a fitness application interface. A central modal window titled "Monday Workout" is open, showing a "Push Day" schedule. It lists four exercises with their respective targets and input fields for sets, reps, and weight. The background interface includes a sidebar with "This Week's Workouts" (Total: 51) and a list of workouts for Monday and Tuesday. The main area shows a "Log Workout" button and a list of logged workouts.

Figure 10.37: Workout Logging Form

Workout Entry Form: Select workout type (Strength, Cardio, Sports, Yoga, Other), enter exercise details including sets, reps, and weight, add cardio duration and intensity, include notes and mood indicator, with automatic calorie calculation.

10.7.4 Progress Analytics

Visualize your fitness progress:




Figure 10.38: Fitness Analytics

Analytics Include: Weight progress chart with target line and goal projection, workout consistency calendar heatmap with streak counter, exercise volume tracking over time, and personal records list.

10.8 Profile & Settings

10.8.1 User Profile

Manage your account information:




Figure 10.39: User Profile Page

Profile Sections:

- Personal Information: Full name, username, email address, phone number, and profile picture upload.
- Account Settings: Change password, email preferences, privacy settings, data export (GDPR), and account deletion.

10.8.2 Telegram Bot Integration

Connect your Telegram account for mobile notifications:




Figure 10.41: Telegram Integration

Integration Steps:

10. Click 'Connect Telegram' button to generate 11-char verification code.
11. Open Telegram and search for @The_Cortana_Project_bot.
12. Send /link followed by the code to complete connection.




Figure 10.42: Telegram Verification Code

Command	Function
/help	Show all available commands
/balance	Check current balance
/budget	View budget status
/add [amount] [category] [desc]	Add expense
/income [amount] [desc]	Add income
/summary	Get weekly financial summary
/news	Get latest news briefing

Table 11.8: Telegram Bot Commands

10.9.3 Dark Mode

Toggle between light and dark themes:




Figure 10.46: Light and Dark Mode Comparison

Dark Mode Benefits: Reduced eye strain in low light, OLED-friendly design for battery savings on mobile, preference saved per device, and automatic theme switching based on system settings.

Chapter 11: Mobile Application - User Interface & Features


11.1 Introduction to the Cortana Mobile App

The Cortana mobile application brings the full power of the web dashboard to smartphones, optimized for on-the-go financial management and lifestyle tracking. Built with Flutter for cross-platform compatibility, the mobile app provides a native-like experience on both iOS and Android devices.

11.2 Mobile App Setup

11.2.2 Account Setup

New User Registration:



The screenshot shows the 'Create Account' screen of the Cortana mobile app. At the top, there is a status bar with the time 10:50 and various icons. Below the status bar is a large, glowing blue circular logo. The title 'Create Account' is centered, followed by the subtitle 'Sign up to get started with Cortana AI'. The registration form consists of five input fields: 'Username' (with a person icon), 'Email' (with an envelope icon), 'Password' (with a lock icon and a toggle for visibility), 'Full Name (Optional)' (with a name tag icon), and 'Phone Number (Optional)' (with a phone icon). A blue 'Sign Up' button is at the bottom of the form.

Figure 11.4: Mobile Registration

Existing User Login:




Figure 11.6: Mobile Login

Login features include username and password fields, 'Remember me' toggle and 'Forgot Password' link.

11.3 Mobile App Home Screen

11.3.1 Dashboard Overview

The mobile home screen presents a financial snapshot at a glance:




Figure 11.7: Mobile Home Dashboard

Dashboard Layout:

- Header Section: Welcome message with user's name, current date, notification bell with badge, and profile picture for settings access.

- Balance Card: Large display of current balance with gradient background, period indicator, and income/expense summary.
- Quick Stats Row: Two side-by-side cards showing Income (green) and Expenses (red) with percentage changes.
- Quick Actions: Four buttons for Add Expense, Add Income, Scan Receipt, and Ask Cortana.
- Recent Transactions: Last 5 transactions with category icon, name, date/time, and color-coded amount.
- Budget Progress Widget: Progress bar with color indicator, spent/total amounts, and days remaining.




Figure 11.8: Balance Card

Quick Actions



Figure 11.9: Quick Action Buttons

Budget Overview




Figure 11.10: Budget Progress Widget

11.3.2 Navigation Bar

Persistent navigation across all screens:




Figure 11.11: Navigation

11.4 Adding Transactions on Mobile

11.4.1 Quick Add Button

The central '+ Add' button provides the fastest way to log transactions:




Figure 11.12: Add Transaction Sheet

Quick Add Flow:

13. Tap Central + Button: Bottom sheet slides up with Expense selected by default.
14. Transaction Type Toggle: Large toggle between Expense and Income with color change.
15. Amount Input: Large number pad with currency display, backspace, and decimal point.
16. Category Selection: Horizontal scrollable category chips with icons.
17. Description: Optional text field with auto-suggestions based on past transactions.
18. Date Picker: Defaults to today with calendar and quick options.
19. Save Button: Validates amount, shows success animation, and closes sheet.



Figure 11.13: Amount Input Pad



Figure 11.14: Category Selection

11.5 Transactions Screen (Mobile)

11.5.1 Transaction List View

Browse all your financial history:




Figure 11.21: Transaction List

Screen Layout:

- Search Bar: Real-time filtering by description, category, amount, or date.
- Filter Chips: Horizontal scrollable options for All, Income, Expenses, This Week, This Month, Custom Range.
- Grouped Transaction List: Entries grouped by date showing category icon, description, date/time, and color-coded amount.
- Swipe Actions: Swipe left to delete (red), swipe right to edit (blue).




Figure 11.22: Swipe Actions

11.5.2 Transaction Details View

Tap any transaction to see full details:




Figure 11.23: Transaction Details

11.7 AI Chat on Mobile

11.7.1 Mobile Chat Interface

Optimized chat experience for mobile:




Figure 11.29: Mobile Chat

Chat Interface:

- Header with Cortana avatar, status, and menu.
- Message list with styled bubbles, timestamps, and typing indicator.
- Input area with attachment, microphone, and send buttons.

11.8 Mobile Budget & Goal Management

11.8.1 Budget Setup Wizard

Easy step-by-step budget creation:




Figure 11.33: Budget Setup Wizard

Setup Steps: Welcome screen with benefits, period selection (Weekly/Monthly/Bi-weekly), budget amount input with AI suggestions, optional category allocation with pie chart, alert preferences, and confirmation summary.

11.8.2 Budget Dashboard

Monitor your budget in real-time:




Figure 11.34: Budget Dashboard

Dashboard Components: Main budget card with circular progress ring, budget status alert banner, category budget cards with mini progress bars, and spending pace indicator with projection.

11.9 Health & Fitness on Mobile

11.9.1 Health Dashboard

Track fitness from your phone:




Figure 12.36: Mobile Health Dashboard

Dashboard Layout: Current weight card with trend and mini chart, BMI indicator with category, workout streak with calendar heatmap, and weekly activity summary.

11.9.2 Log Weight and Workout

11:07 100% 6g

← Log Weight

Weight (kg)

70.5

Body Fat % (Optional)

15.5

Weigh-In Date

2026-01-26

Notes (Optional)

Morning weight, after workout

Log Weight

Figure 11.37: Weight Logging

Weight Entry: Simple interface with number pad, unit toggle, date picker, optional notes, and progress photo option.

11:08 100% 60%

← Log Workout



Exercise Name

e.g., Bench Press

Sets (Optional) Reps (Optional)

3 10

Weight (kg) Duration (min)

60.0 45

Notes (Optional)

Felt strong today, increased weight

Log Workout

Figure 11.38: Workout Logging

Workout Logging: Select workout type, enter exercises with sets/ reps/ weight, timer between sets, add cardio details, and complete with summary and mood selector.

11.11 Profile & Settings (Mobile)

11.11.1 Mobile Profile Screen




Figure 11.41: Mobile Profile

Profile Layout: Profile header with photo, name, username, email. Account settings, app settings (currency, language, notifications, theme, biometric). Connected services status. Support links and logout.

Chapter 12: Advanced AI Features Deep Dive

12.1 Natural Language Processing

12.1.1 Understanding Conversational Input

Cortana's NLP engine understands various ways of expressing the same transaction:




Figure 12.1: NLP Input Flexibility

Examples of Equivalent Inputs for '\$25 lunch at Chipotle':


- 'I spent \$25 on lunch at Chipotle'
- 'Lunch was 25 dollars at Chipotle'
- 'Paid 25 bucks for Chipotle'
- 'Chipotle lunch \$25'
- 'Twenty-five dollars Chipotle food'

NLP Capabilities:


Capability	Description	Examples
Entity Extraction	Identifies amounts, categories, merchants, dates	\$25, yesterday, last Tuesday
Intent Recognition	Determines user's goal	'spent', 'paid' = Add expense
Context Awareness	Remembers conversation history	Handles follow-up questions

Table 12.1: NLP Capabilities


12.2 AI Chat with Context Retrieval



12.3 Workout Plant Generation



13.2 Telegram Bot Intergration



General Conclusion and Future Recommendations

Conclusion

Cortana AI Assistant successfully demonstrates the practical application of advanced AI technologies to personal productivity management. The system's RAG architecture, combining vector-based semantic search with large language models, enables contextually aware assistance that references user-specific historical data. This represents a significant advancement over both traditional productivity applications and generic AI chatbots.

The two-month AI research phase yielded a robust RAG implementation using FAISS for vector storage and semantic search. Automatic vectorization converts every financial transaction to embeddings without user intervention, enabling natural language queries across entire transaction histories. The three-tier AI fallback system (Groq → Gemini → Ollama) achieved 100% availability, ensuring continuous service despite individual provider limitations.

Rigorous testing validated system performance: 94% accuracy in natural language expense parsing, sub-millisecond vector search performance, 100% AI availability, and sub-200ms API response times. User testing demonstrated 67% reduction in expense logging time, 89% satisfaction with AI insights, and 4.4/5 overall satisfaction rating.

The project makes several technical contributions: Personal Data RAG (first known implementation for personal finance), Three-Tier AI Fallback (novel reliability architecture), and Lebanese Localization (dual-currency support, Arabic language, local news sources).

Future Recommendations

While Cortana successfully addresses critical needs in personal productivity management, several areas offer opportunities for future enhancement:

Voice Command Integration: Implement native voice interfaces for hands-free operation on web and mobile platforms.

Receipt OCR: Add automatic expense extraction from receipt photos using computer vision.

Real-time Collaboration: Enable multi-user workspaces for shared budgets and family expense tracking.

iOS Application: Extend mobile support to iOS platform for broader user reach.

Bank Integration: Connect to financial institutions for automatic transaction import.

Predictive Analytics: Implement ML-based expense forecasting and budget recommendations.

Advanced Health Features: Integrate with fitness trackers and add nutrition tracking capabilities.

In conclusion, Cortana represents a significant advancement in AI-powered personal productivity. By continually refining and expanding its functionalities based on these recommendations, the platform can further enhance its effectiveness, scalability, and impact on users' daily lives.

References

- [1] P. Lewis, E. Perez, A. Piktus, et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," Proceedings of NeurIPS, pp. 9459-9474, 2020.
- [2] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks," Proceedings of EMNLP, pp. 3982-3992, 2019.
- [3] A. Vaswani, N. Shazeer, N. Parmar, et al., "Attention Is All You Need," Proceedings of NeurIPS, pp. 5998-6008, 2017.
- [4] J. Wei, X. Wang, D. Schuurmans, et al., "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models," Proceedings of NeurIPS, pp. 24824-24837, 2022.
- [5] J. Johnson, M. Douze, and H. Jégou, "Billion-scale similarity search with GPUs," IEEE Transactions on Big Data, vol. 7, no. 3, pp. 535-547, 2019.
- [6] Meta AI, "FAISS: A Library for Efficient Similarity Search," <https://github.com/facebookresearch/faiss>, accessed Jan. 2025.
- [7] S. Whittaker, L. Terveen, and B. A. Nardi, "Let's stop pushing the envelope and start addressing it," Human-Computer Interaction, vol. 15, pp. 75-106, 2011.
- [8] H. Shum, X. He, and D. Li, "From Eliza to XiaoIce: challenges and opportunities with social chatbots," Frontiers of IT & EE, vol. 19, pp. 10-26, 2018.
- [9] Y. Zhang et al., "DIALOGPT: Large-Scale Generative Pre-training for Conversational Response Generation," Proceedings of ACL, pp. 270-278, 2020.
- [10] S. Russell and P. Norvig, Artificial Intelligence: A Modern Approach, 4th ed., Pearson, 2020.
- [11] FastAPI, "FastAPI Framework," <https://fastapi.tiangolo.com/>, accessed Jan. 2025.
- [12] PostgreSQL Global Development Group, "PostgreSQL Documentation," <https://www.postgresql.org/docs/>, accessed Jan. 2025.
- [13] React Team, "React Documentation," <https://react.dev/>, accessed Jan. 2025.

- [14] Flutter Team, "Flutter Documentation," <https://flutter.dev/docs>, accessed Jan. 2025.
- [15] Groq, "Groq LPU Inference Engine," <https://groq.com/>, accessed Jan. 2025.
- [16] Google AI, "Gemini API," <https://ai.google.dev/>, accessed Jan. 2025.
- [17] Ollama, "Get up and running with large language models locally," <https://ollama.ai/>, accessed Jan. 2025.
- [18] Hugging Face, "Sentence Transformers Documentation," <https://www.sbert.net/>, accessed Jan. 2025.
- [19] OpenAI, "ChatGPT," <https://openai.com/chatgpt>, accessed Jan. 2025.
- [20] Apple Inc., "Siri," <https://www.apple.com/siri/>, accessed Jan. 2025.